

Neural Collaborative Filtering

1.Introduction

2.Improvement

- Inner Product --> Neural network prediction


3.Core

- Generalized Matrix Factorization(GMF)
- Neural Network(MLP)
- Fusion

4.Result

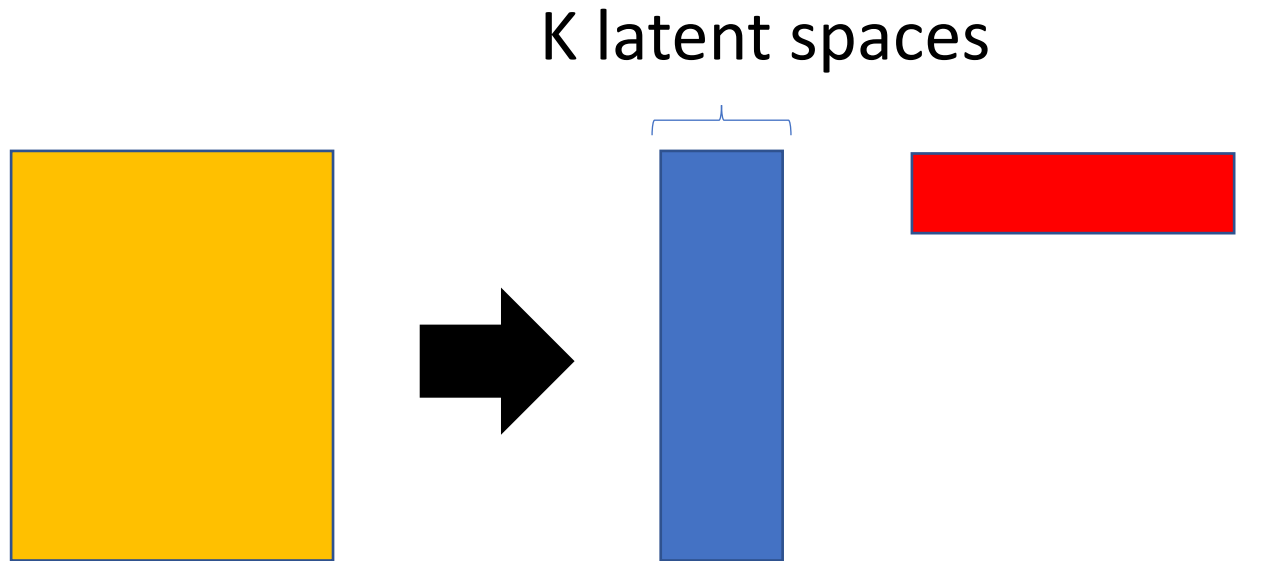
Introduction

- NCF is an improvement for normal interaction method like MF.
- It uses the deep neural network(DNN) to learn the interaction function.
- Combines Linear(GMF) and non-Linear(MLP) model.



Improvement

Matrix Factorization decompose a $M \times N$ user-item interaction matrix into a $M \times K$ user matrix and a $K \times N$ item matrix. And it is described that **we have projected the users and the items to the same K latent spaces.**



If a user U and an item I are alike in the latent space, user U is prone to buy item I . And we use the inner product to judge whether the user vector is close to the item vector, that's where using inner product as the possibility of interaction between a user and an item comes from


Improvement

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

← items →

↑ users

(a) user-item matrix



(b) user latent space

Using simple method like inner product or cosine index will lead to a great ranking loss

But the latent space method has a limitation

$$s_{23}(0.66) > s_{12}(0.5) > s_{13}(0.4)$$

$$s_{41}(0.6) > s_{43}(0.4) > s_{42}(0.2),$$

We cannot lay U_4 correctly

Improvement

Normal Matrix Factorization:

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i = \sum_{k=1}^K p_{uk} q_{ik},$$

Use the inner production of user vector \mathbf{U} and item vector \mathbf{I} to predict the interaction between the user and the item.

NCF:

$$\hat{y}_{ui} = \text{NCF}(p_{uk} q_{ik})$$

Use Neural network to calculate the interaction between user \mathbf{U} and item \mathbf{I} on given input user vector \mathbf{U} and item vector \mathbf{I}

Core

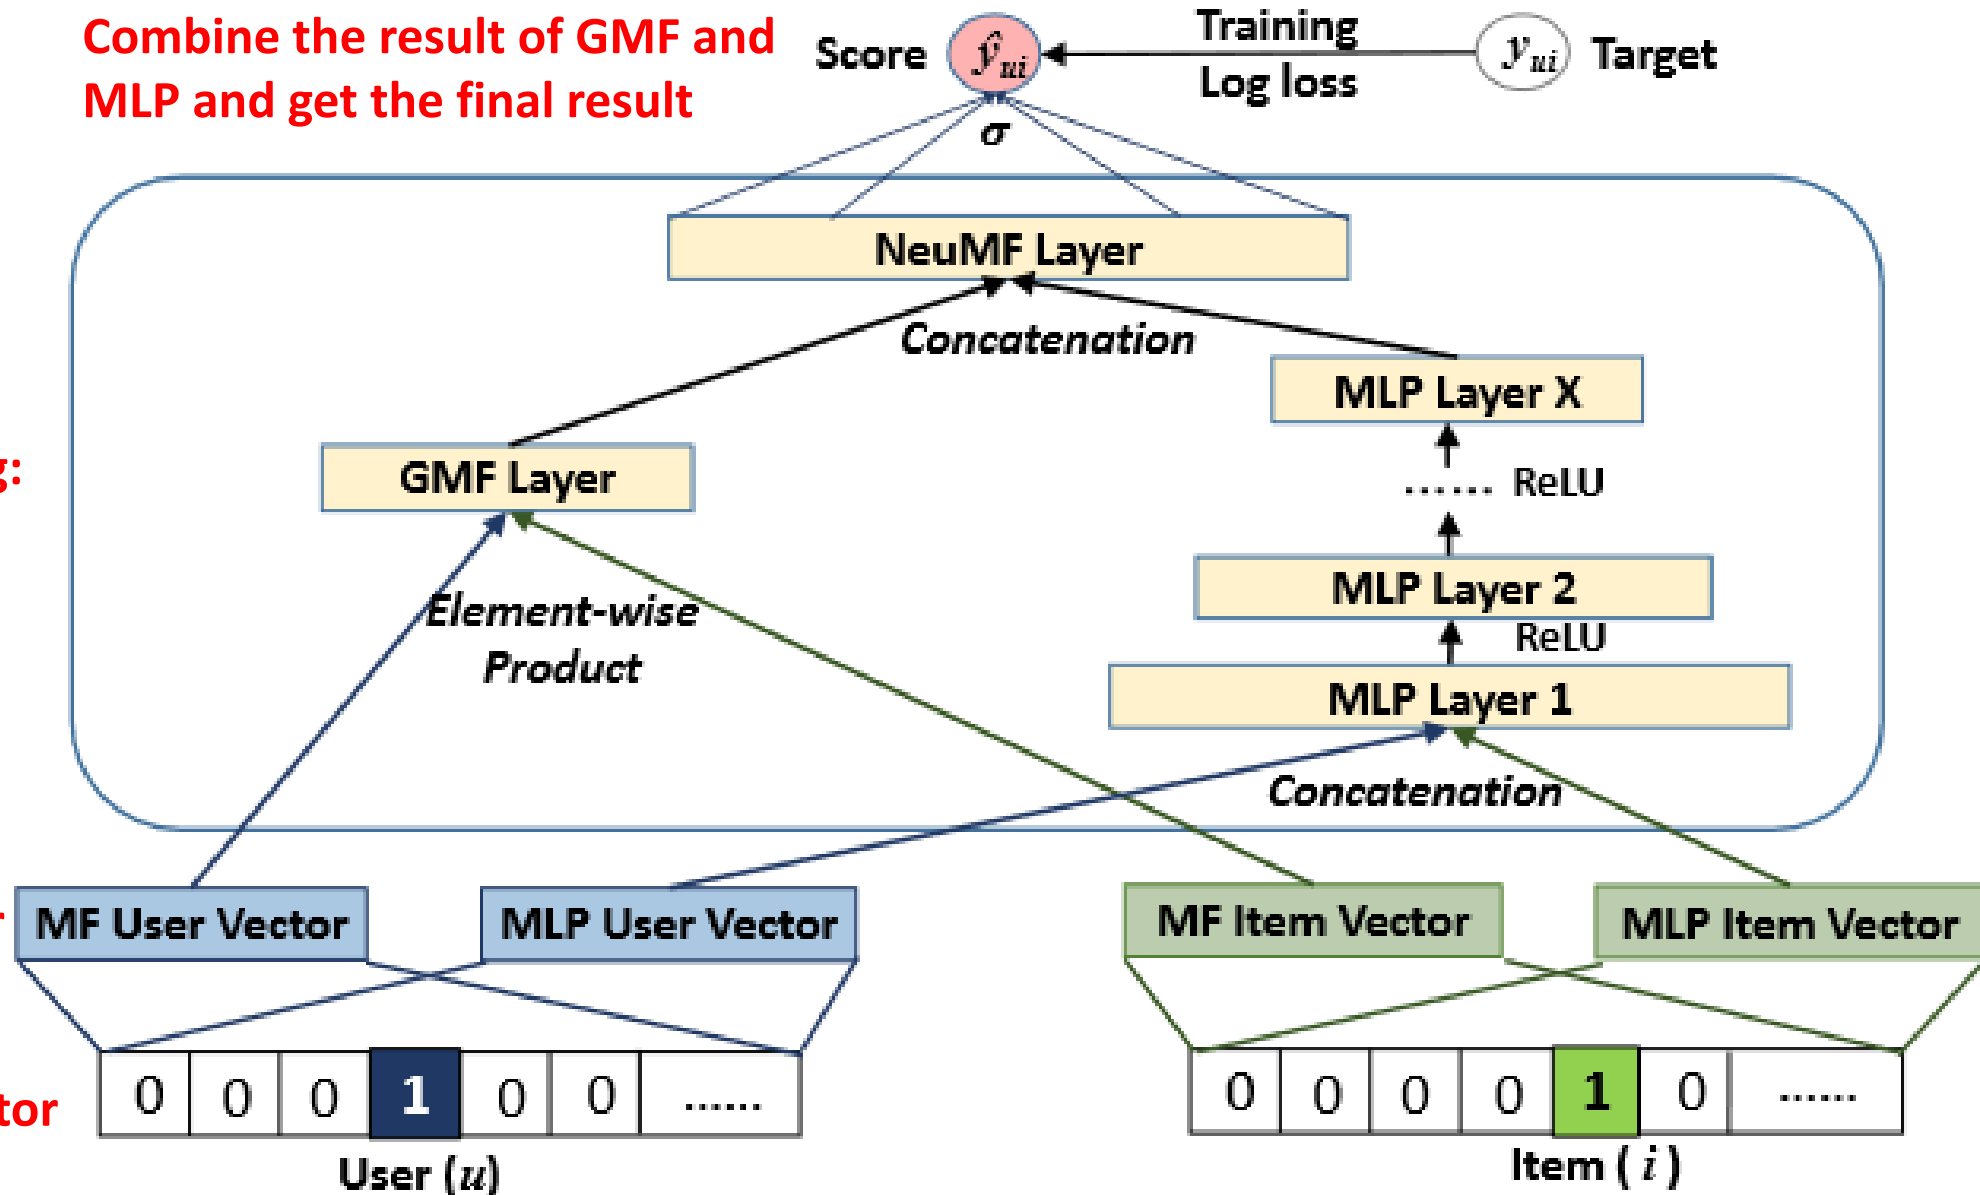
Combine the result of GMF and MLP and get the final result



Predict the result using:
GMF/MLP

User/Item latent vector

User/Item One-hot vector



Core

$$\hat{y}_{u,i} = NCF(U_i^{GMF}, I_j^{GMF}, U_i^{MLP}, I_j^{MLP}, MLP - Net, h^{GMF}, h^{MLP}, \alpha)$$

$\hat{y}_{u,i}$ The predicted interaction between user u and item i (output)

$MLP - Net$ The MLP neural network

h^{GMF} A parameter for GMF part


h^{MLP} A parameter for MLP part

α Voting weight for the fusion of two parts

Core

Generalized Matrix Factorization(Linear Part):

$$\phi^{GMF} = \mathbf{p}_u^G \odot \mathbf{q}_i^G$$




(inner product: the sum of each element in that vector, 18)

Core

Multi-Layer Perceptron(non-Linear Part):

$$\phi^{MLP} = a_L(\mathbf{W}_L^T(a_{L-1}(\dots a_2(\mathbf{W}_2^T \begin{bmatrix} \mathbf{p}_u^M \\ \mathbf{q}_i^M \end{bmatrix} + \mathbf{b}_2)\dots)) + \mathbf{b}_L)$$




Core

Combines the result of both GMF and MLP using different weight(α)

$$\hat{y}_{ui} = \sigma(\mathbf{h}^T \begin{bmatrix} \phi^{GMF} \\ \phi^{MLP} \end{bmatrix})$$

$$\mathbf{h} \leftarrow \begin{bmatrix} \alpha \mathbf{h}^{GMF} \\ (1 - \alpha) \mathbf{h}^{MLP} \end{bmatrix}$$

$$\hat{y}_{ui} = \text{Sigmoid}(\text{sum}(\))$$



Core

GMF,MLP: pre-training using Adam

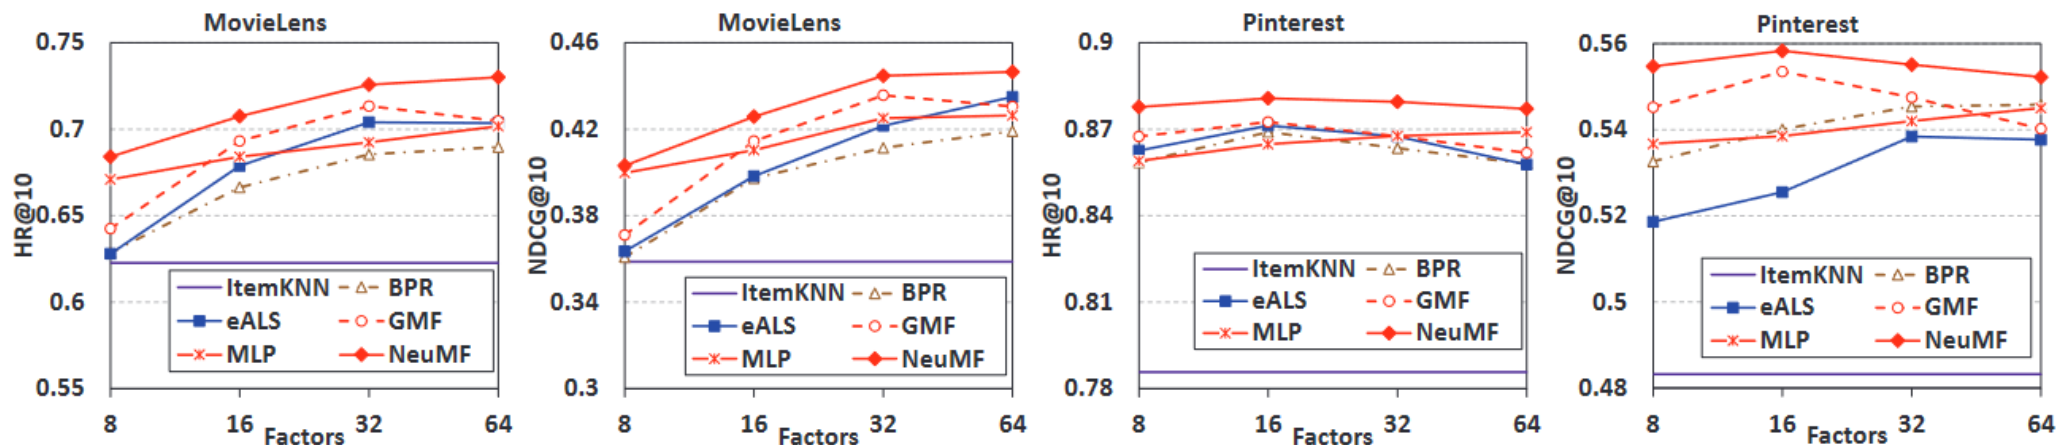
$$\phi^{GMF} = \mathbf{p}_u^G \odot \mathbf{q}_i^G$$

$$\phi^{MLP} = a_L(\mathbf{W}_L^T(a_{L-1}(\dots a_2(\mathbf{W}_2^T \begin{bmatrix} \mathbf{p}_u^M \\ \mathbf{q}_i^M \end{bmatrix} + \mathbf{b}_2)\dots)) + \mathbf{b}_L)$$

Use the trained GMF and MLP, then only tweak \mathbf{h} ,
using vanilla SGD, based on cross-entropy loss


$$\mathbf{h} \leftarrow \begin{bmatrix} \alpha \mathbf{h}^{GMF} \\ (1 - \alpha) \mathbf{h}^{MLP} \end{bmatrix} \quad \begin{aligned} L &= - \sum_{(u,i) \in \mathcal{Y}} \log \hat{y}_{ui} - \sum_{(u,j) \in \mathcal{Y}^-} \log(1 - \hat{y}_{uj}) \\ &= - \sum_{(u,i) \in \mathcal{Y} \cup \mathcal{Y}^-} y_{ui} \log \hat{y}_{ui} + (1 - y_{ui}) \log(1 - \hat{y}_{ui}). \end{aligned}$$

Result



(a) MovieLens — HR@10 (b) MovieLens — NDCG@10 (c) Pinterest — HR@10 (d) Pinterest — NDCG@10

Figure 4: Performance of HR@10 and NDCG@10 *w.r.t.* the number of predictive factors on the two datasets.



(a) MovieLens — HR@K (b) MovieLens — NDCG@K (c) Pinterest — HR@K (d) Pinterest — NDCG@K

Figure 5: Evaluation of Top-K item recommendation where K ranges from 1 to 10 on the two datasets.


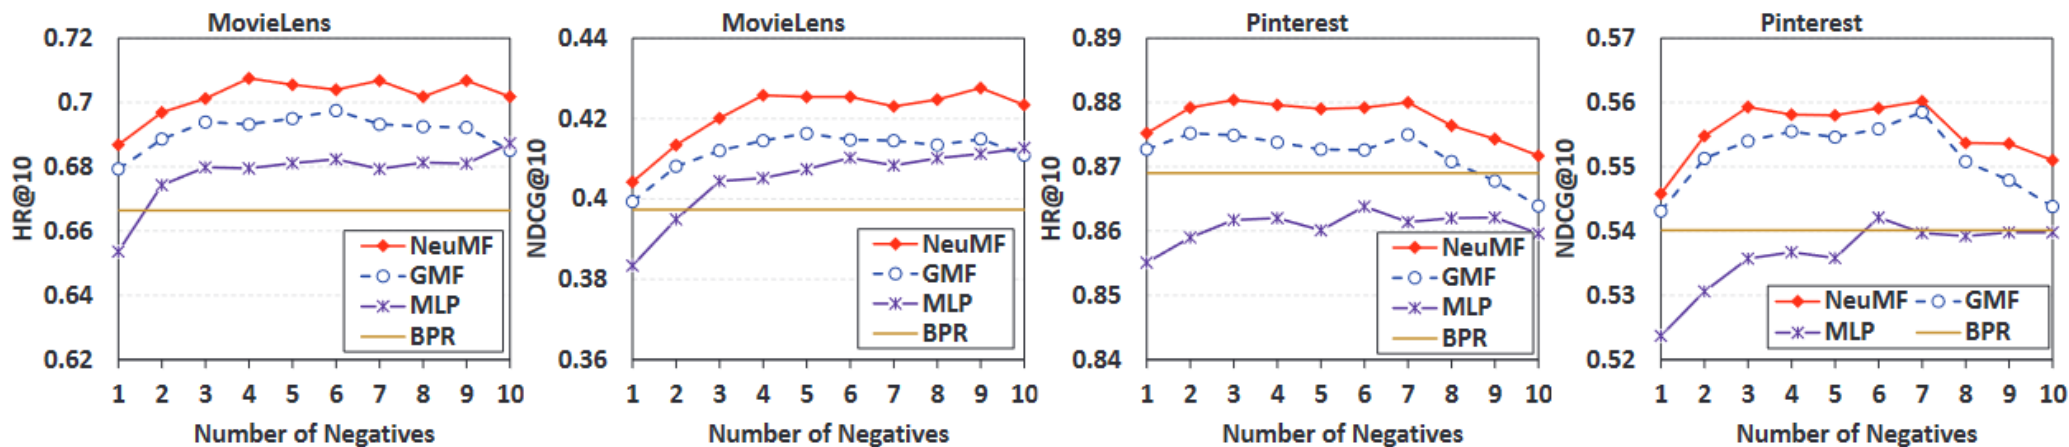


Figure 6: Training loss and recommendation performance of NCF methods *w.r.t.* the number of iterations on MovieLens (factors=8).



(a) MovieLens — HR@10 (b) MovieLens — NDCG@10 (c) Pinterest — HR@10 (d) Pinterest — NDCG@10

Figure 7: Performance of NCF methods *w.r.t.* the number of negative samples per positive instance (factors=16). The performance of BPR is also shown, which samples only one negative instance to pair with a positive instance for learning.